



Capstone Project

Edge computing device programming for AI projects

Lecture 3

Dr. Wilton Fok &
Carol Chen

Recap: Training classification model

	Operation	Results
Dataset preparation	Preparing the catdog classification dataset	Get a prepared dataset
Train model	<pre>\$ cd ~/jetson-inference/python/training/classification \$ python3 train.py data/cat_dog --model-dir=models/cat_dog --batch-size=1 --workers=1 --epochs=1</pre>	Obtaining a .pth model
Convert model	<pre>\$ python3 onnx_export.py --model-dir=models/cat_dog</pre>	Obtaining a .onnx model
Run model	<pre>/home/nvidia/jetson-inference/build/aarch64/bin/imagenet -- model= models/resnet18.onnx --input_blob=input_0 -- output_blob=output_0 --labels data/cat_dog/labels.txt data/cat_dog/test/cat</pre>	Visualize the output of the model

Training with your own dataset

A group of dataset


Two important parameter should be defined

- Where to store your data (***YourDataPath***)
- Where to store your model (***YourModelPath***)


Following the similar steps to finish model training, conversion and inference.

Step1: Dataset preparation

- Now you are familiar with how to train classification model.
- Before training, we should collect the dataset, and organize them similar to the catdog dataset.



```
--data/catdog
-----train
-----cat
-----dog
-----test
-----cat
-----dog
-----validation
-----cat
-----dog
-----labels.txt
```



Step1: Dataset preparation

- Now, you can download the images you want to classify from google.
- Remember the image among train, validation and test is around 7:2:1
- After finishing the dataset, run “check_dataset.py” to verify whether the dataset structure is valid.

```
nvidia@nvidia-desktop: ~/jetson-inf...training/classification/data/cat_dog - □ ×
File Edit Tabs Help
nvidia@nvidia-desktop:~/jetson-inference/python/training/classification/data/cat_dog$ python3 check_directory.py
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/labels.txt directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/train directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/test directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/val directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/train/cat directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/test/cat directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/val/cat directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/train/dog directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/test/dog directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/val/dog directory does exist.
nvidia@nvidia-desktop:~/jetson-inference/python/training/classification/data/cat_dog$
```

Correct directory

```
nvidia@nvidia-desktop: ~/jetson-Infer...n/training/classification/data/cat_dog - □ ×
File Edit Tabs Help
nvidia@nvidia-desktop:~/jetson-inference/python/training/classification/data/cat_dog$ python3 check_classification_dataset.py
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/labels.txt directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/train directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/test directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/val directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/train/cat directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/test/cat directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/val/cat directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/train/dog directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/test/dog directory does exist.
/home/nvidia/jetson-inference/python/training/classification/data/cat_dog/val/dog directory does exist.
nvidia@nvidia-desktop:~/jetson-inference/python/training/classification/data/cat_dog$
```

Missing "val/cat"

```

1  import os
2
3  data_path = 'data/cat_dog'
4
5  labels = os.path.join(data_path, 'labels.txt')
6
7  Train_dir = os.path.join(data_path, 'train')
8  Test_dir = os.path.join(data_path, 'test')
9  Val_dir = os.path.join(data_path, 'val')
10
11  checkTrain = os.path.isdir(Train_dir)
12  checkTest = os.path.isdir(Test_dir)
13  checkVal = os.path.isdir(Val_dir)
14
15  if labels:
16      print(labels, 'does exist')
17  else:
18      print("Missing labels.txt")
19
20  if checkTrain:
21      print(Train_dir, " does exist")
22      findcat = os.path.isdir(os.path.join(Train_dir, 'cat'))
23      finddog = os.path.isdir(os.path.join(Train_dir, 'dog'))
24      if findcat:
25          print(os.path.join(Train_dir, 'cat'), " does exist")
26      else:
27          print("Missing train/cat directory!")
28      if finddog:
29          print(os.path.join(Train_dir, 'dog'), " does exist")
30      else:
31          print("Missing train/dog directory!")
32  else:
33      print("Missing train directory!")
34
35  if checkTest:
36      print(Test_dir, " does exist")
37      findcat = os.path.isdir(os.path.join(Test_dir, 'cat'))
38      finddog = os.path.isdir(os.path.join(Test_dir, 'dog'))
39      if findcat:
40          print(os.path.join(Test_dir, 'cat'), " does exist")
41      else:
42          print("Missing test/cat directory!")
43      if finddog:
44          print(os.path.join(Test_dir, 'dog'), " does exist")
45      else:
46          print("Missing test/dog directory!")
47  else:
48      print("Missing test directory!")
49
50  if checkVal:
51      print(Val_dir, " does exist")
52      findcat = os.path.isdir(os.path.join(Val_dir, 'cat'))
53      finddog = os.path.isdir(os.path.join(Val_dir, 'dog'))
54      if findcat:
55          print(os.path.join(Val_dir, 'cat'), " does exist")
56      else:
57          print("Missing val/cat directory!")
58      if finddog:
59          print(os.path.join(Val_dir, 'dog'), " does exist")
60      else:
61          print("Missing val/dog directory!")
62  else:
63      print("Missing val directory!")
64

```

Step 2: Training

- After getting the dataset, similarly, we need to open a terminal, and:

```
$ cd ~/jetson-inference/python/training/classification
```

```
$ python3 train.py YourDataPath --model-dir=YourModelPath --batch-size=1 --workers=1 --epochs=10
```

Then the model will begin to train.

*Also, Please keep the path of **YourModelPath** and **YourDataPath** in mind! They are necessary in the following*

Some concepts of commands and AI

During this time, I will introduce what's the meaning of the commands entered.

Remember: Terminal is a controller communicating with your computer!

Commands translation

(1) **`cd ~/jetson-inference/python/training/classification`**

Translation: Entering the base location, because our AI files are written and we will begin here.

(2) **`python3 train.py data/cat_dog --model-dir=models/cat_dog --batch-size=1 --workers=1 --epochs=1`**

Translation: Train a model. We use the data stored in data/cat_dog to train and the model will be stored in models/cat_dog. The batch size, workers and epochs are set to be 1.


(3) **`python3 onnx_export.py --model-dir=models/cat_dog`**

Translation: We convert the model in models/cat_dog to onnx.

(4) **`/home/nvidia/jetson-inference/build/aarch64/bin/imagenet --model=models/resnet18.onnx --input_blob=input_0 --output_blob=output_0 --labels data/cat_dog/labels.txt data/cat_dog/test/cat`**

Translation: We use the ".../imagenet" to execute running images, with loading the model

- Training with your own data




After finishing your training, you would find this two files in "mytraining"

- Converting the Model to ONNX

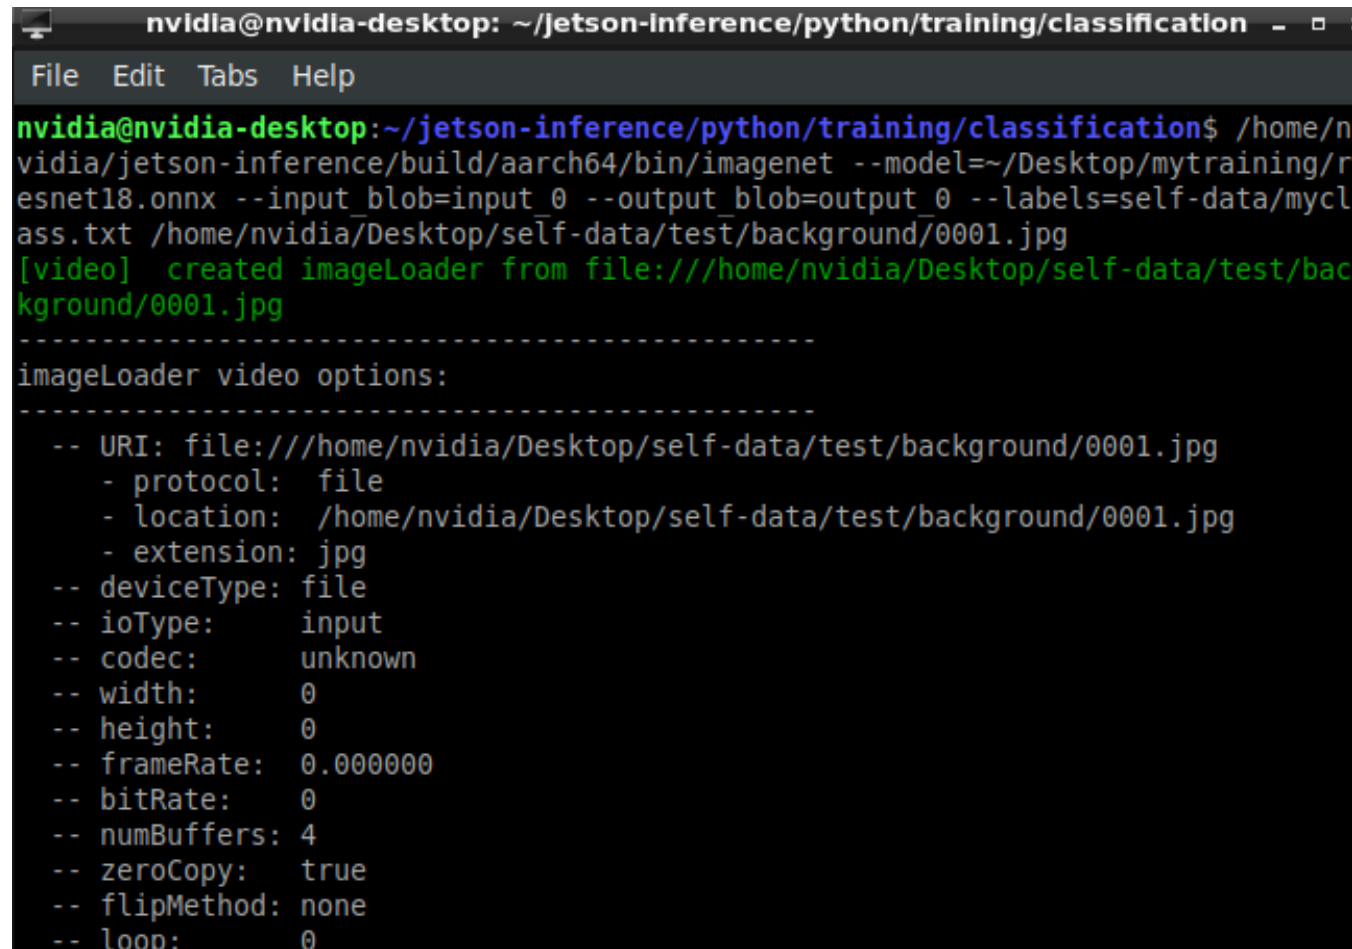
This will create a model called resnet18.onnx under:

YourModelPath




- Processing Images with TensorRT

```
$ /home/nvidia/jetson-inference/build/aarch64/bin/imagenet --  
model=YourModelPath/resnet18.onnx --input_blob=input_0 --output_blob=output_0 --  
labels=self-data/class.txt yourOwenPicturePath
```



```
nvidia@nvidia-desktop: ~/jetson-inference/python/training/classification  
File Edit Tabs Help  
nvidia@nvidia-desktop:~/jetson-inference/python/training/classification$ /home/n  
vidia/jetson-inference/build/aarch64/bin/imagenet --model=~/.Desktop/mytraining/r  
esnet18.onnx --input_blob=input_0 --output_blob=output_0 --labels=self-data/mycl  
ass.txt /home/nvidia/Desktop/self-data/test/background/0001.jpg  
[video] created imageLoader from file:///home/nvidia/Desktop/self-data/test/bac  
kground/0001.jpg  
-----  
imageLoader video options:  
-----  
-- URI: file:///home/nvidia/Desktop/self-data/test/background/0001.jpg  
  - protocol: file  
  - location: /home/nvidia/Desktop/self-data/test/background/0001.jpg  
  - extension: jpg  
-- deviceType: file  
-- ioType: input  
-- codec: unknown  
-- width: 0  
-- height: 0  
-- frameRate: 0.000000  
-- bitRate: 0  
-- numBuffers: 4  
-- zeroCopy: true  
-- flipMethod: none  
-- loop: 0
```

Task2: Object detection training




Object Detection is a computer vision task in which the goal is to detect and locate objects of interest in an image or video. The task involves identifying the position and boundaries of objects in an image, and classifying the objects into different categories.

Classification



CAT

Object Detection



CAT, DOG, DUCK

Classification




CAT

Classification + Localization



CAT

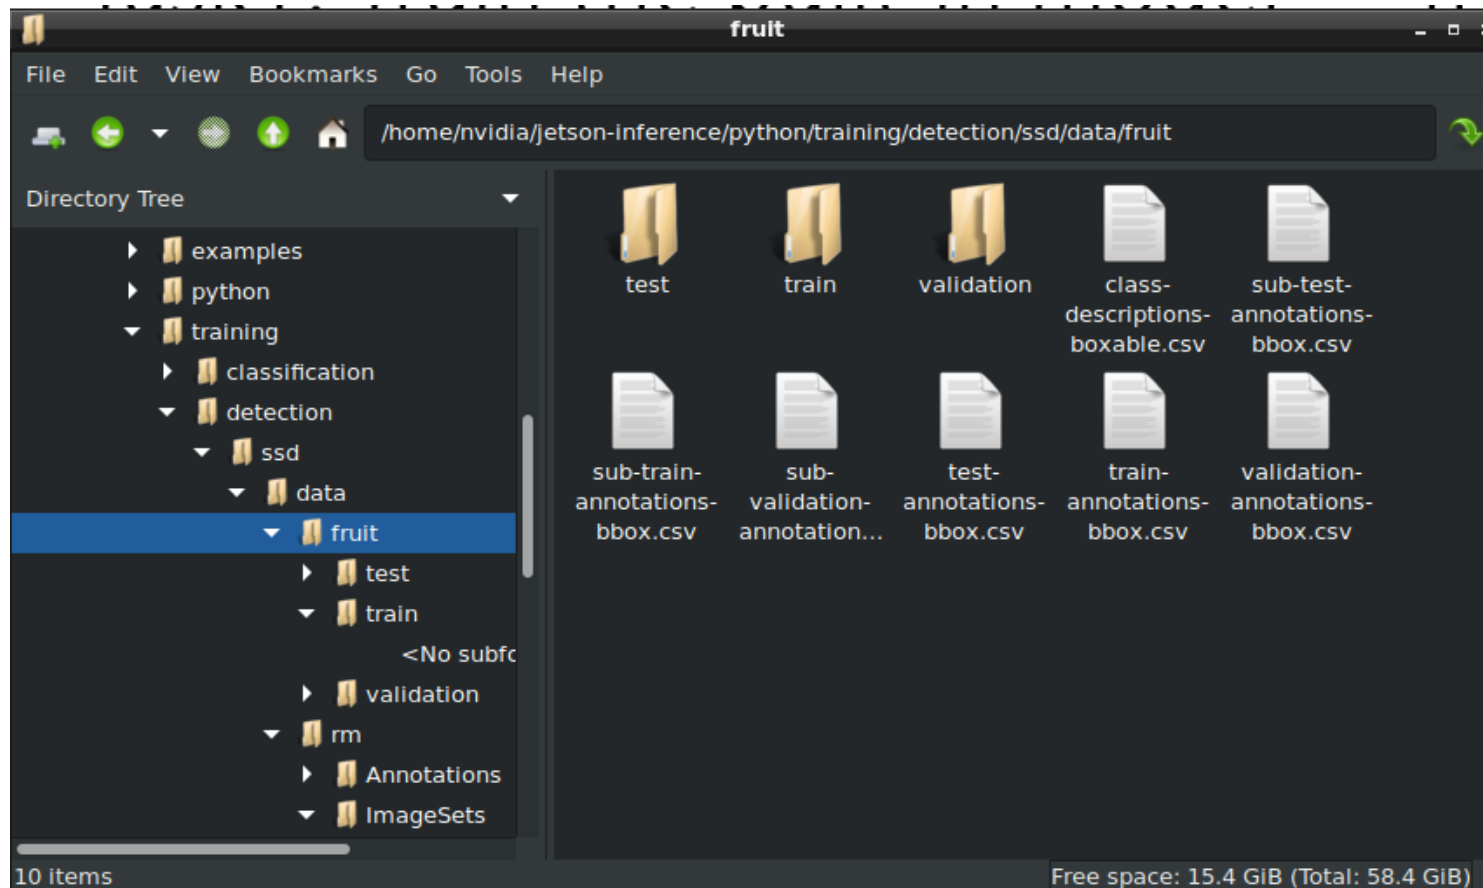
Object Detection





CAT, DOG, DUCK

Train the built-in model --- fruit

- Mission: Train a fruit recognition model using these images, and use it to detection different kinds of fruit.



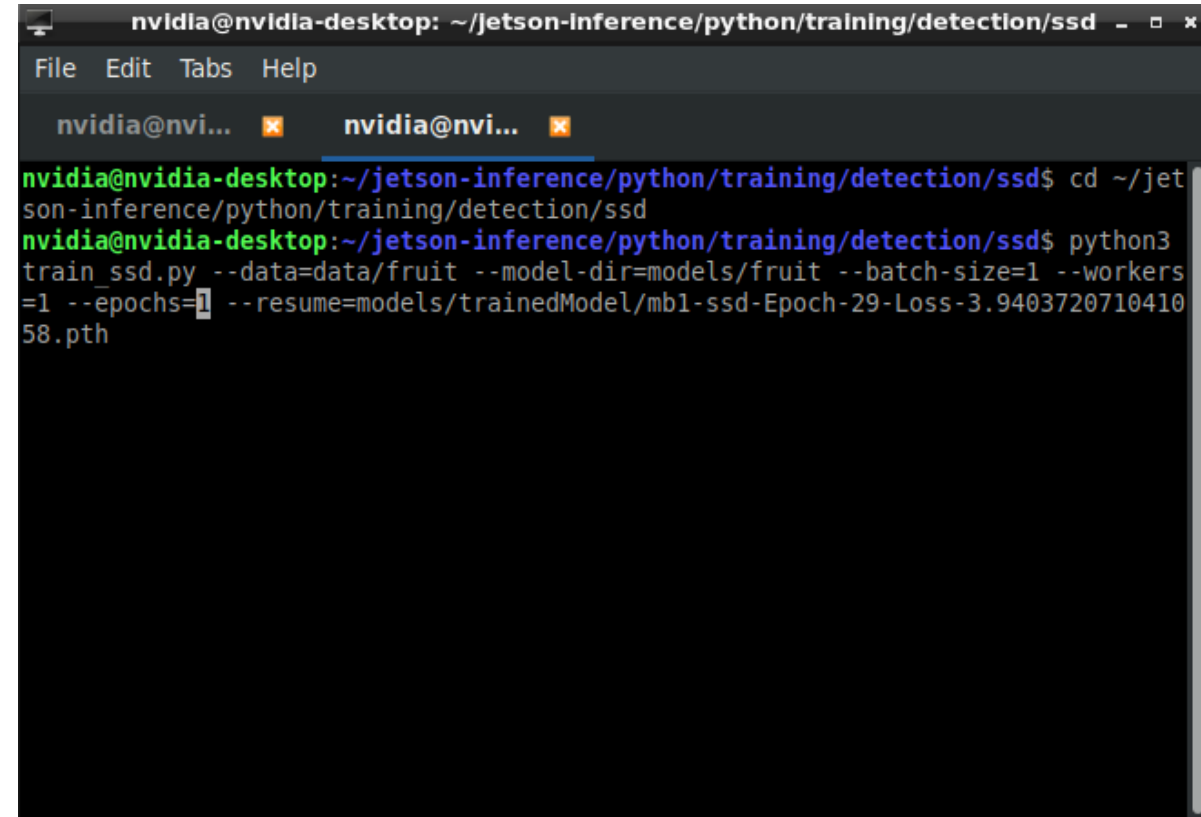
The fruit dataset have been stored in */home/nvidia/jetson-inference/python/training/detection/ssd/data* path.



Step 1: Train model

- To train the model, we need firstly download the pretrained model [mb1-ssd-Epoch-29-Loss-3.940372071041058.pth](#) and put it in models/trainedModel
- Download the dataset [fruit](#) and put it in data/fruitt
- then open a terminal:

```
$ cd ~/jetson-  
inference/python/training/detection/ssd  
$ python3 train_ssd.py --data=data/fruit --  
model-dir=models/fruit --batch-size=1 --  
workers=1 --epochs=1 --  
resume=models/trainedModel/mb1-ssd-  
Epoch29-Loss-3.940372071041058.pth
```



The screenshot shows a terminal window titled "nvidia@nvidia-desktop: ~/jetson-inference/python/training/detection/ssd". The terminal displays the following commands and output:

```
nvidia@nvidia-desktop: ~/jetson-inference/python/training/detection/ssd$ cd ~/jetson-inference/python/training/detection/ssd  
nvidia@nvidia-desktop: ~/jetson-inference/python/training/detection/ssd$ python3 train_ssd.py --data=data/fruit --model-dir=models/fruit --batch-size=1 --workers=1 --epochs=1 --resume=models/trainedModel/mb1-ssd-Epoch-29-Loss-3.940372071041058.pth
```

Step 1: Train model

- During training, the terminal will show lots of information.

2020-07-10 13:14:12 - Epoch: 0, Step: 10/1287, Avg Loss: 12.4240, **Avg Regression Loss** 3.5747, **Avg Classification Loss**: 8.8493

2020-07-10 13:14:12 - Epoch: 0, Step: 20/1287, Avg Loss: 9.6947, Avg Regression Loss 4.1911, Avg Classification Loss: 5.5036

2020-07-10 13:14:13 - Epoch: 0, Step: 30/1287, Avg Loss: 8.7409, Avg Regression Loss 3.4078, Avg Classification Loss: 5.3332

2020-07-10 13:14:13 - Epoch: 0, Step: 40/1287, Avg Loss: 7.3736, Avg Regression Loss 2.5356, Avg Classification Loss: 4.8379

2020-07-10 13:14:14 - Epoch: 0, Step: 50/1287, Avg Loss: 6.3461, Avg Regression Loss 2.2286, Avg Classification Loss: 4.1175


...

2020-07-10 13:19:26 - Epoch: 0, Validation Loss: 5.6730, Validation Regression Loss 1.7096, Validation Classification Loss: 3.9634

2020-07-10 13:19:26 - Saved model models/fruit/mb1-ssd-Epoch-0-Loss-5.672993580500285.pth


Avg Regression Loss:

- Avg Regression Loss is the average error of the bounding box locations.




Avg Classification Loss

- Avg Classification Loss: Avg Classification Loss is the average error of the object classifications.
- E.g. **Dog** **Cat** **Duck** (0% loss)



Step 2: Convert model

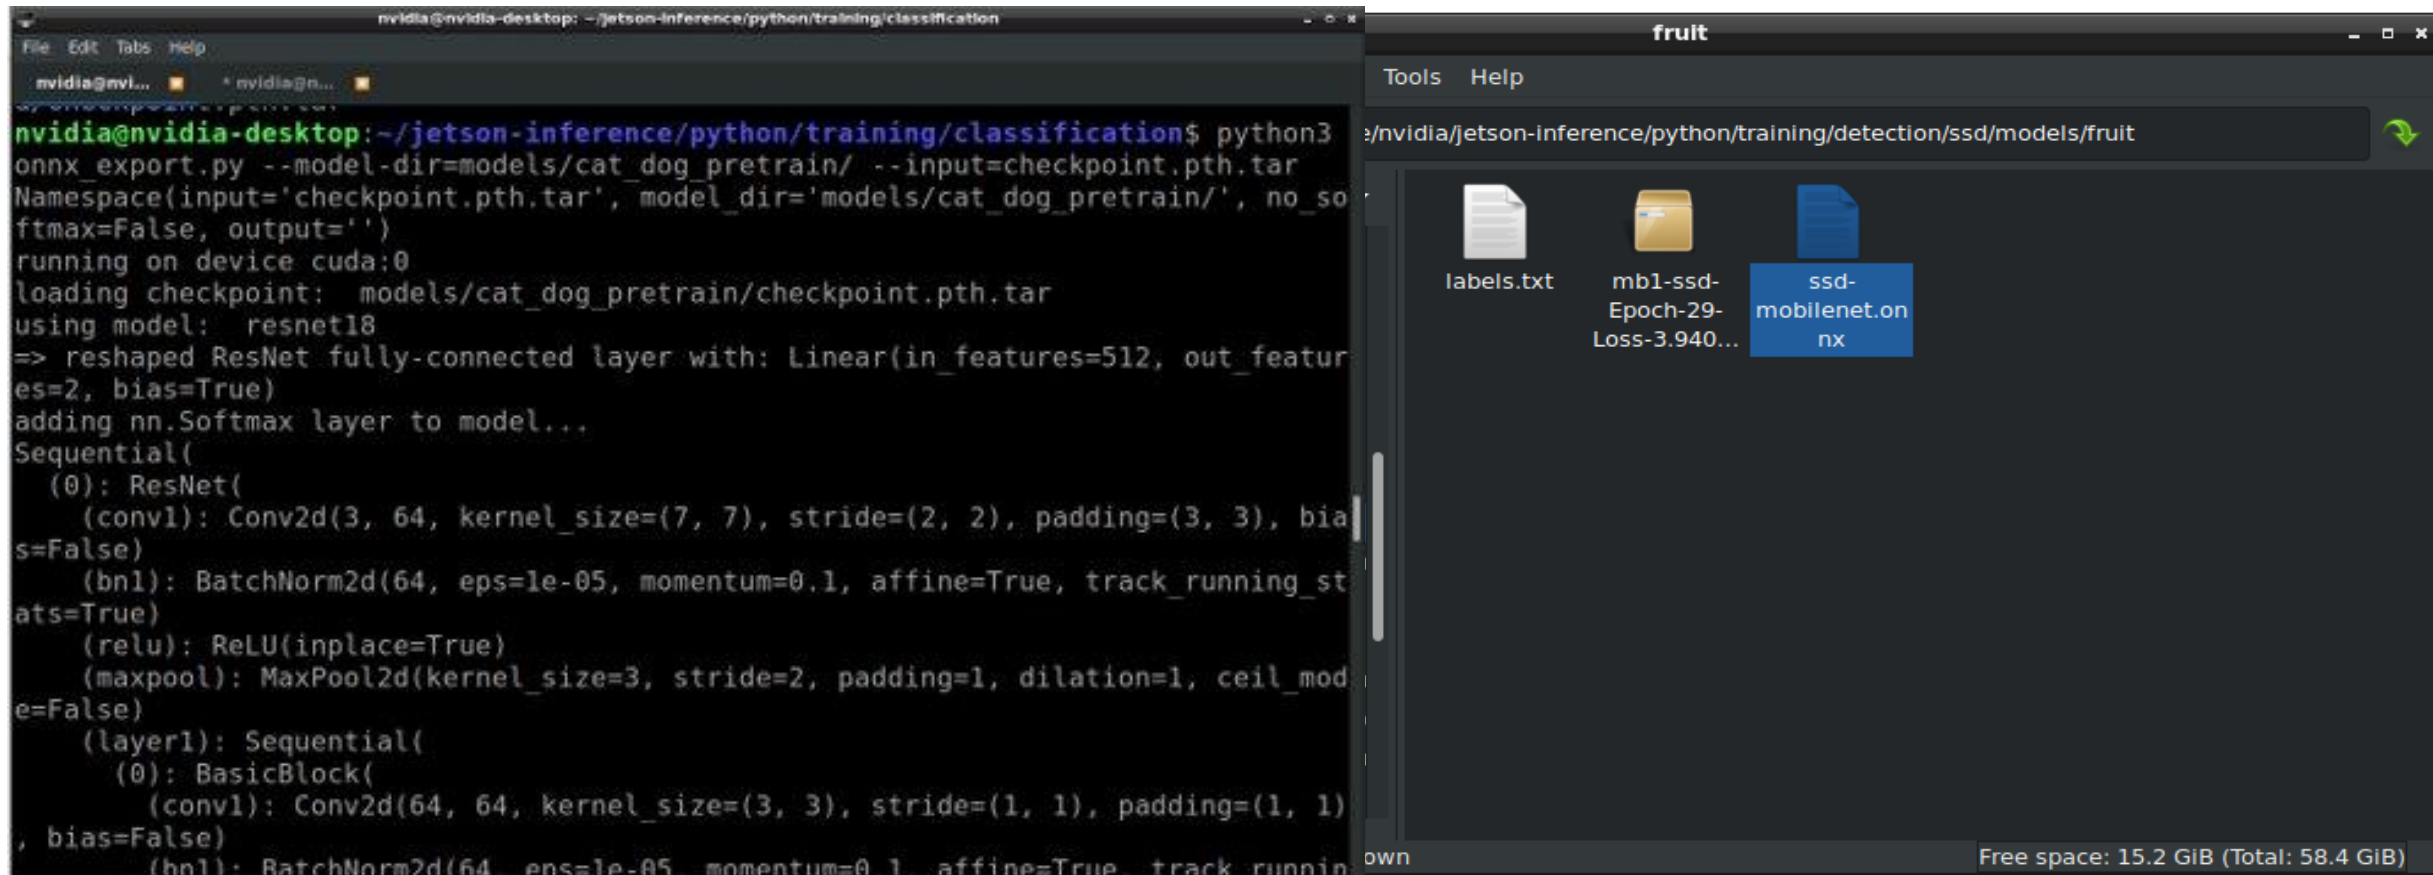
After training, the model folder will be:



Then, we need to convert the model using:

```
$ python3 onnx_export.py --  
input=models/fruit/mb1-ssd-Epoch-29-Loss-  
3.940372071041058.pth --model-dir=models/fruit/
```

When you find a file with .onnx in the folder, it succeeds!



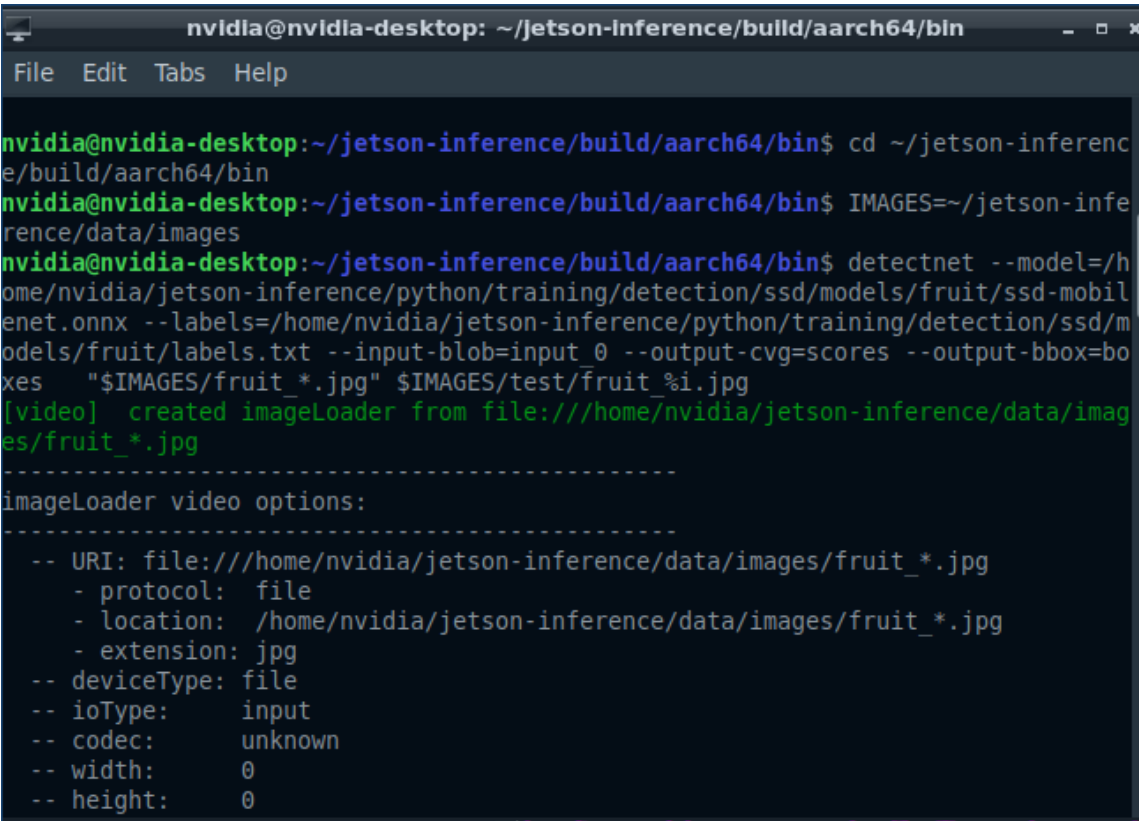
Step 3: Inference model with image

- Now you get a model runnable! In the next step, we need to use it to make prediction on different images.
- Take fruit for example:

```
$ cd ~/jetson-inference/build/aarch64/bin
```

```
$ IMAGES=/home/nvidia/jetson-  
inference/data/images
```

```
$ detectnet --model=/home/nvidia/jetson-  
inference/python/training/detection/ssd/models/fru  
it/ssd-mobilenet.onnx --labels=/home/nvidia/jetson-  
inference/python/training/detection/ssd/models/fru  
it/labels.txt --input-blob=input_0 --output-cvg=scores  
--output-bbox=boxes "$IMAGES/fruit_*.jpg"  
$IMAGES/test/fruit_%i.jpg
```




```
nvidia@nvidia-desktop: ~/jetson-inference/build/aarch64/bin
File Edit Tabs Help

nvidia@nvidia-desktop:~/jetson-inference/build/aarch64/bin$ cd ~/jetson-inference/build/aarch64/bin
nvidia@nvidia-desktop:~/jetson-inference/build/aarch64/bin$ IMAGES=/home/nvidia/jetson-inference/data/images
nvidia@nvidia-desktop:~/jetson-inference/build/aarch64/bin$ detectnet --model=/home/nvidia/jetson-inference/python/training/detection/ssd/models/fruit/ssd-mobilenet.onnx --labels=/home/nvidia/jetson-inference/python/training/detection/ssd/models/fruit/labels.txt --input-blob=input_0 --output-cvg=scores --output-bbox=boxes "$IMAGES/fruit_*.jpg" $IMAGES/test/fruit_%i.jpg
[video] created imageLoader from file:///home/nvidia/jetson-inference/data/images/fruit_*.jpg
-----
imageLoader video options:
-----
-- URI: file:///home/nvidia/jetson-inference/data/images/fruit_*.jpg
- protocol: file
- location: /home/nvidia/jetson-inference/data/images/fruit_*.jpg
- extension: jpg
-- deviceType: file
-- ioType: input
-- codec: unknown
-- width: 0
-- height: 0
```

Step 3: Inference model with image

- You can find the result at */home/nvidia/jetson-inference/data/images/test*



Fruit detection with live camera

```
$ cd ~/jetson-inference/build/aarch64/bin
```

```
$ detectnet --model=/home/nvidia/jetson-
```

```
inference/python/training/detection/ssd/models/fruit/ssd-mobilenet.onnx --
```

```
labels=/home/nvidia/jetson-
```

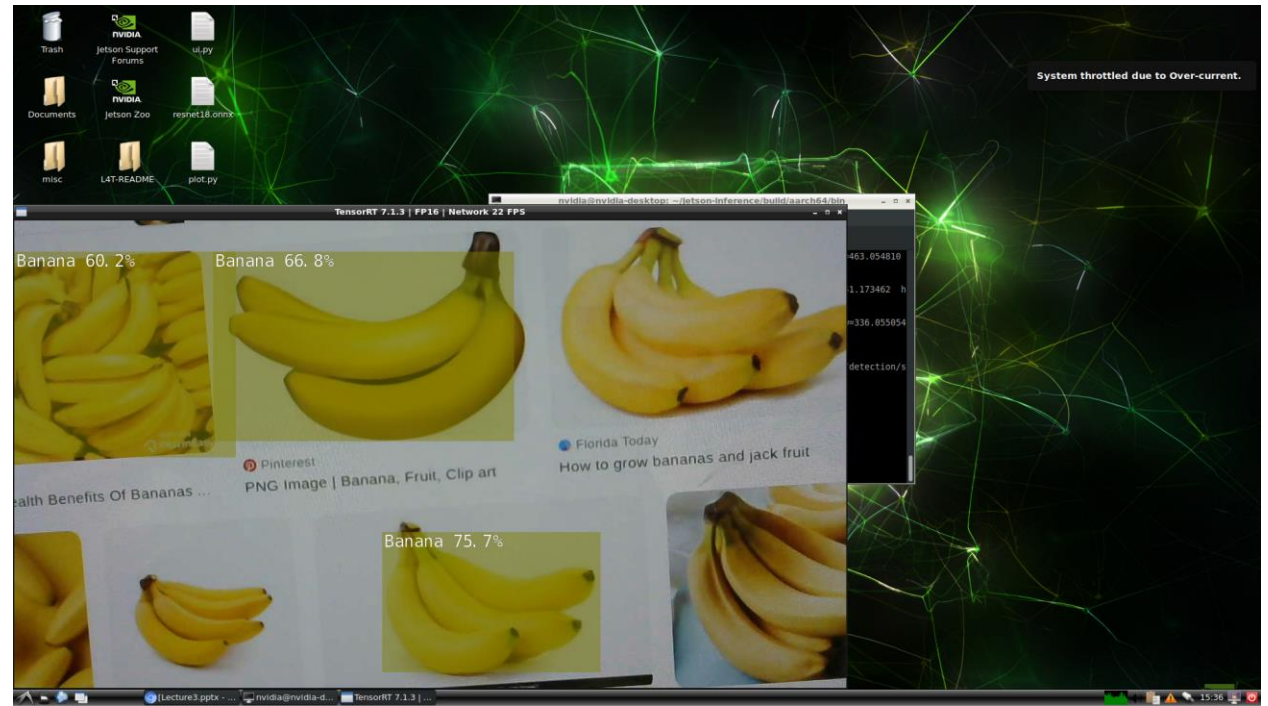
```
inference/python/training/detection/ssd/models/fruit/labels.txt --input-blob=input_0 --
```

```
output-cvg=scores --output-bbox=boxes /dev/video0
```

```
nvidia@nvidia-desktop: ~/jetson-Inference/build/aarch64/bin
File Edit Tabs Help

nvidia@nvi... x nvidia@nvi... x

nvidia@nvidia-desktop:~/jetson-inference/build/aarch64/bin$ cd ~/jetson-inference/build/aarch64/bin
nvidia@nvidia-desktop:~/jetson-inference/build/aarch64/bin$ detectnet --model=/home/nvidia/jetson-inference/python/training/detection/ssd/models/fruit/ssd-mobilenet.onnx --labels=/home/nvidia/jetson-inference/python/training/detection/ssd/models/fruit/labels.txt --input-blob=input_0 --output-cvg=scores --output-bbox=boxes /dev/video0
```



Coding Your Own Object Detection Program

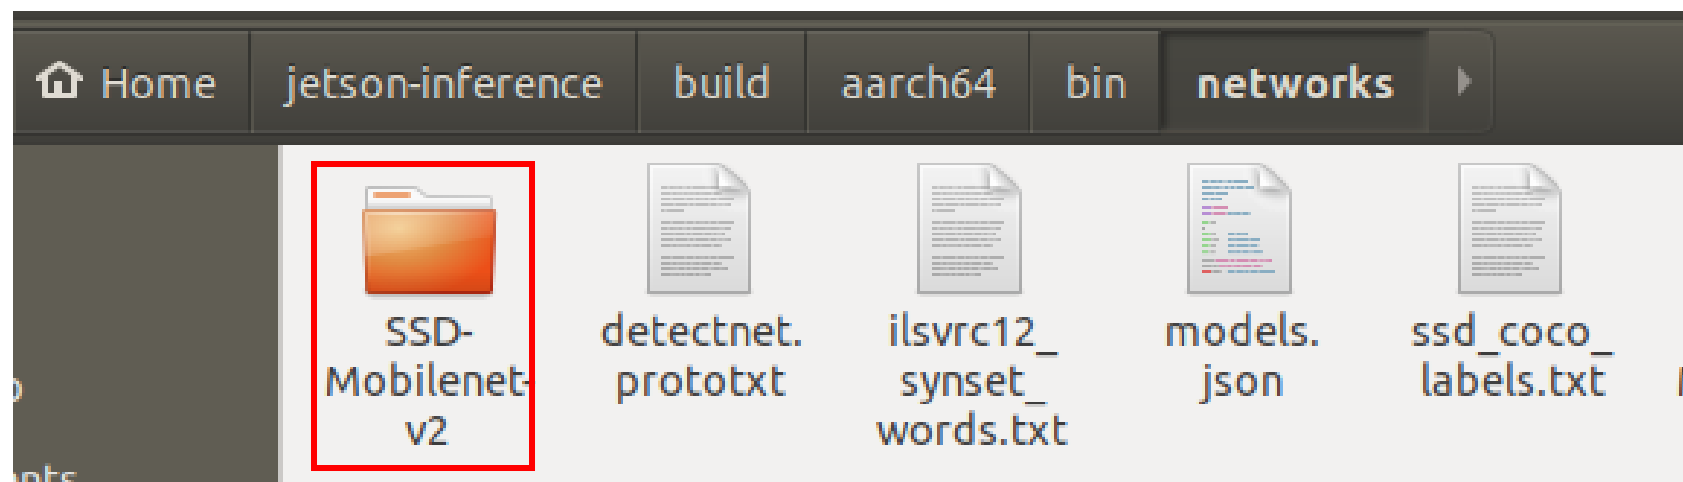
First, open up your text editor of choice and create a new file named **my-detection.py**. Below we'll assume that you'll save it on your host device under **home/jetson-inference/examples/**,

At the top of the source file, we'll import the Python modules that we're going to use in the script. Add import statements to load the [jetson_inference](#) and [jetson_utils](#) modules used for object detection and camera capture.

```
=====
#from jetson_inference import detectNet
#from jetson_utils import videoSource, videoOutput
import jetson.inference
import jetson.utils
```


Loading the Detection Model

Next download the [SSD-Mobilenet-v2 model](#) and put the folder in the path: **/home/nvidia/jetson-inference/build/aarch64/bin/networks**



Then use the following line to create a detectNet object instance that loads the 91-class SSD-Mobilenet-v2 model:

```
=====
net = detectNet("ssd-mobilenet-v2", threshold=0.5)
```

Opening the Camera Stream

To connect to the camera device for streaming, we'll create an instance of the videoSource object:

```
=====
camera = videoSource("/dev/video0")    # '/dev/video0' for V4L2
```

Display Loop

Next, we'll create a video output interface with the [videoOutput](#) object and create a main loop that will run until the user exits:

```
=====
display = videoOutput("display://0") # 'my_video.mp4' for file

while display.IsStreaming(): # main loop will go here
```

Camera Capture

The first thing that happens in the main loop is to capture the next video frame from the camera. `camera.Capture()` will wait until the next frame has been sent from the camera and loaded into GPU memory.

```
=====
display = videoOutput("display://0") # 'my_video.mp4' for file

while display.IsStreaming(): # main loop will go here
    img = camera.Capture()

    if img is None: # capture timeout
        continue
```


Next the detection network processes the image with the `net.Detect()` function. It takes in the image from `camera.Capture()` and returns a list of detections:

=====

```
#from jetson_inference import detectNet
#from jetson_utils import videoSource, videoOutput
import jetson.inference
import jetson.utils

net = detectNet("ssd-mobilenet-v2", threshold=0.5)
camera = videoSource("/dev/video0")    # '/dev/video0' for V4L2
display = videoOutput("display://0") # 'my_video.mp4' for file
while display.IsStreaming():
    img = camera.Capture()
    if img is None: # capture timeout
        continue
    detections = net.Detect(img)
```

You can add
a `print(detections)` statement here,
and the coordinates,
confidence, and class info will
be printed out to the terminal
for each detection result.



Rendering

Finally we'll visualize the results with OpenGL and update the title of the window to display the current performance:

=====

```
while display.IsStreaming():
```

```
    img = camera.Capture()
```

```
    if img is None: # capture timeout
```

```
        continue
```

```
    detections = net.Detect(img)
```

```
    display.Render(img)
```

```
    display.SetStatus("Object Detection | Network {:.0f} FPS".format(net.GetNetworkFPS()))
```

Running the Program

```
$ python3 my-detection.py
```

example

```
-- ClassID: 5
-- Confidence: 0.927246
-- Left: 2.44141
-- Top: 60.791
-- Right: 499
-- Bottom: 274.475
-- Width: 496.559
-- Height: 213.684
-- Area: 106107
-- Center: (250.721, 167.633)
```

```
ge] loaded '/home/nvid
5, 3 channels)
DetectNet.Detection object
-- ClassID: 5
-- Confidence: 0.927246
-- Left: 2.44141
-- Top: 60.791
-- Right: 499
-- Bottom: 274.475
-- Width: 496.559
-- Height: 213.684
-- Area: 106107
-- Center: (250.721, 16
```

Some concepts

- **ClassID:** The ClassID represents the category or class of the detected object.
- **Confidence:** The Confidence score indicates how certain the model is about its prediction for the detected object.
- These values define the coordinates of the bounding box around the detected object.
 - Left & Top:** Left and Top specify the coordinates of the top-left corner of the bounding box.
 - Right & Bottom:** Right and Bottom specify the coordinates of the bottom-right corner. These coordinates are essential for locating the object in the image.
- **Width:** Width and Height are derived from the bounding box coordinates and represent the dimensions of the bounding box. They show how wide and tall the detected object is in the image.
- **Height:** Width and Height are derived from the bounding box coordinates and represent the dimensions of the bounding box. They show how wide and tall the detected object is in the image.
- **Area:** The Area is the total size of the bounding box, calculated as $\text{Width} \times \text{Height}$.
- **Center:** The Center represents the midpoint of the bounding box, calculated using the coordinates.

Assignment 3



Find out the coordinates, confidence, and class info of two image detection results including ClassID, Confidence, Left, Top, Right, Bottom, Width, Height, Area, Center.



Please upload the code (your-detection.py) to github, a doc or pdf to the moodle (including github link, the detected image and the detection results)



Deadline: 03/03/2025 11:59pm

Tips

- Change the video input to one of the image input
- Choose one of the class output
- Learn to use debug function in Visual Studio Code